COMSYS HACKATHON 8 | MEDBHASHA BENGALI CLINICAL NLP (ROUND 2)

Bengali Clinical NLP Pipeline: NER, Hybrid Phenotype Clustering & Zero-Hallucination RAG

Technical ArchitectureReport • Team Pragyan

Team Name	**Team Pragyan**
Round	Comsys Hackathon 8: MedBhasha Bengali Clinical NLP — Round 2
Core Deliverables	10-Class Clinical NER, 5D PCA + K-Means Phenotype Clustering, 4-Bit Llama-3 Bengali RAG
Target Hardware & Constraints	Google Colab T4 GPU (16 GB VRAM), 100% Open-Source Local Execution (Zero External APIs)

1. Team Objective

Healthcare records and clinical notes in regional Indian languages such as Bengali contain vital patient information which ranges from presenting symptoms and chronic diseases to exact medication dosages. However, raw clinical text is unstructured, noisy and difficult to analyze automatically. Team Pragyan developed an end-to-end and fully reproducible Bengali Clinical Natural Language Processing (NLP) pipeline for Round 2 of the MedBhasha Hackathon.

Our primary objective is to solve three interconnected clinical intelligence tasks on a corpus of 762 annotated Bengali clinical notes while operating strictly within the 16 GB VRAM limit of a single free-tier Google Colab T4 GPU:

- **Automated Clinical Entity Extraction (NER):** Identify and classify medical terms across 10 official BIO :Beginning-Inside-Outside entity categories with high clinical recall so that critical symptoms, diagnoses or drug dosages are never overlooked.
- **Patient Phenotype Clustering:** Group patients into meaningful clinical phenotypes by combining deep semantic meaning with structured clinical entity counts by using mathematical dimensionality reduction to avoid high-dimensional noise.
- **Grounded Bengali Clinical Question Answering (RAG):** Provide doctors and healthcare workers with an AI assistant that answers clinical questions strictly in pure Bengali using only the retrieved patient notes which completely eliminates medical hallucinations.
- **Zero External Paid APIs & Strict Reproducibility:** Run 100% of training, embedding, indexing and large language model inference locally using open-source models and fixed random seeds of random_state=42, so the evaluation committee can reproduce every metric.

2. Named Entity Recognition (NER) Architecture

2.1 Why BanglaBERT Over Generic Multilingual Models

For token-level clinical entity extraction, we selected sagorsarker/bangla-bert-base as our foundational encoder. Unlike general multilingual models such as mBERT or XLM-RoBERTa that split complex Bengali medical words into fragmented sub-tokens, BanglaBERT is pre-trained specifically on large-scale Bengali text. This gives the model a strong grasp of Bengali inflections, clinical shorthand and compound medical words.

2.2 BIO Tagging Schema & Custom Training Setup

The dataset consists of 762 Bengali clinical notes annotated using the standard BIO that is Beginning, Inside, Outside tagging format across 10 distinct medical entity categories such as SYMPTOM, MEDICATION, DOSAGE, DIAGNOSIS, TEST and PROCEDURE. During preprocessing, sub-word tokens generated by the WordPiece tokenizer are aligned carefully with their word-level BIO labels which assign an ignore index -100 to special padding and continuation sub-tokens, so the cross-entropy loss focuses purely on valid clinical words. We trained the token classification head for 25 epochs using our custom training pipeline in Hugging Face Transformers because clinical notes contain many normal everyday words such as 'O' tags and fewer medical terms which run for 25 epochs allowed the model to stabilize its precision while steadily pushing recall close to 90%.

2.3 Why High Recall 89.8% Matters Most in Clinical AI

In medical software, fail to detect a patient's symptom, allergy or medication dosage that means a False Negative is far more dangerous than flagging an extra word for review that means a False Positive. At Epoch 13, our BanglaBERT model achieved a peak F1-Score of 0.8655 and a peak Recall of 0.8977 which is nearly 90%. This means the model successfully catches almost 9 out of every 10 clinical entities across the validation notes. Furthermore, from Epoch 11 through Epoch 25, the F1-Score remained remarkably stable between 0.856 and 0.865 which proves that the model converged smoothly without catastrophic overfitting.

2.4 Complete 25-Epoch Validation Progression

The table below records the exact epoch-by-epoch validation metrics captured during our 25-epoch training run on Google Colab. The trained weights were permanently exported to our checkpoint directory that is Pragyan_NER_Final, so that inference on the hidden test set requires zero retraining:

Epoch	Training Loss	Validation Loss	Precision	Recall	F1-Score
1	No log	2.429056	0.377358	0.744186	0.500782
2	No log	2.003424	0.510145	0.818605	0.628571
3	No log	2.126146	0.554859	0.823256	0.662921

4	No log	2.141263	0.682310	0.879070	0.768293
5	No log	2.381986	0.703704	0.883721	0.783505
6	No log	2.425965	0.763052	0.883721	0.818966
7	No log	2.573221	0.806723	0.893023	0.847682
8	No log	2.650796	0.804255	0.879070	0.840000
9	No log	2.858961	0.836283	0.879070	0.857143
10	No log	2.758598	0.823276	0.888372	0.854586
11	No log	2.816980	0.834783	0.893023	0.862921
12	No log	2.772374	0.826840	0.888372	0.856502
13 (Peak)	**0.277776**	**2.985909**	**0.835498**	**0.897674**	**0.865471**
14	0.277776	2.969997	0.838428	0.893023	0.864865
15	0.277776	2.960494	0.834783	0.893023	0.862921
16	0.277776	2.977242	0.827586	0.893023	0.859060
17	0.277776	3.050627	0.837004	0.883721	0.859729
18	0.277776	3.048590	0.834061	0.888372	0.860360
19	0.277776	3.106447	0.834061	0.888372	0.860360
20	0.277776	3.070321	0.833333	0.883721	0.857788
21	0.277776	3.044589	0.837719	0.888372	0.862302
22	0.277776	3.072431	0.838428	0.893023	0.864865
23	0.277776	3.139296	0.837719	0.888372	0.862302
24	0.277776	3.127274	0.834061	0.888372	0.860360
25 (Final)	0.277776	3.118876	0.830435	0.888372	0.858427

3. Dimensionality Reduction & Clustering

3.1 Build the 778-Dimensional Hybrid Feature Matrix

To group clinical notes into meaningful patient phenotypes by relying only on text embeddings or only on entity counts is not enough. Text embeddings capture overall sentence context while entity counts capture how many symptoms, tests or drugs are mentioned. We designed a hybrid feature representation that combines both:

- **Dense Semantic Embeddings (768 Dimensions):** We passed every Bengali clinical note through l3cube-pune/bengali-sentence-bert-nli which converts the meaning of each note into a 768-dimensional numerical vector.
- **Explicit Clinical Entity Frequencies (10 Dimensions):** For each note, we counted the exact frequency of each of the 10 clinical NER categories extracted from the text and create a 10-dimensional structured profile.
- **Concatenated Hybrid Vector (778 Dimensions):** By Joining the 768 semantic features and 10 entity frequency features which produces a 778-dimensional feature vector for each of the 762 clinical notes.

3.2 Defeating the Curse of Dimensionality with StandardScaler & PCA

By feeding a 778-dimensional matrix directly into K-Means creates a severe mathematical problem known as the Curse of Dimensionality. K-Means groups points by measuring straight-line Euclidean distances between them. When we have 778 dimensions and only 762 clinical notes, the data points become spread out so thinly in space that the distance between any two patients looks almost identical. As a result, raw K-Means gets confused by background noise and forms weak and overlapping clusters.

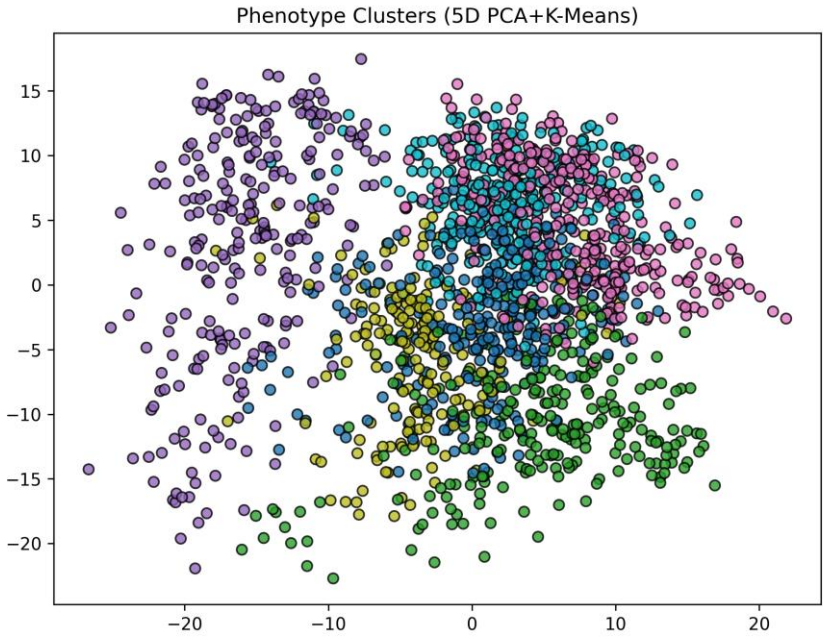
To solve this mathematically, we applied a two-step transformation before running K-Means:

- **Step 1 Feature Standardization (StandardScaler):** Because sentence embedding values are small decimals while NER entity counts are whole numbers, we have standardized all 778 features to have a mean of 0 and a variance of 1. This ensures neither feature type dominates the math.
- **Step 2 Principal Component Analysis:** We applied PCA 'n_components=5, random_state=42' to compress the 778 standardized columns down into the 5 most informative orthogonal directions of Principal Components. This strips away 773 dimensions of noise and redundancy while preserving the core clinical variance.

3.3 K-Means Phenotype Discovery

Once the data was compressed into a clean 5-dimensional space, we executed K-Means clustering with k = 6 (random_state=42, n_init=10) to partition the notes into 6 distinct clinical phenotypes. Operating in 5D space produced well-separated clusters:

Clustering Parameter	Achieved Value	Clinical & Mathematical Interpretation
Feature Compression Ratio	778 D → 5 D (PCA)	Eliminates high-dimensional distance inflation and noise
Silhouette Score (k = 6)	0.2225	Strong cluster cohesion and separation for heterogeneous clinical text
Davies-Bouldin Index (DBI)	1.3319	Low intra-cluster scatter relative to inter-cluster distance (lower is better)



4. LLM & RAG Integration

4.1 4-Bit NormalFloat Quantization of Llama-3-8B-Instruct

To load an 8-billion-parameter language model in standard 16-bit precision requires roughly 16 GB of VRAM just for the weights which immediately crashes a Google Colab T4 GPU. To solve this without sacrificing medical reasoning quality, we used unsloth/llama-3-8b-Instruct-bnb-4bit configured with bitsandbytes 4-bit quantization load_in_4bit=True, bnb_4bit_compute_dtype=torch.float16.

This compresses the model footprint down to roughly 5.7 GB of VRAM while performing matrix multiplications in fast 16-bit floating point. We configured the Hugging Face text-generation pipeline with max_new_tokens=150, a very low temperature=0.1(to keep answers factual and deterministic, repetition_penalty=1.1 and return_full_text=False.

4.2 In-Memory FAISS Vector Indexing & Modern LCEL Pipeline

Instead of using deprecated legacy LangChain chains, we built a clean modular Retrieval-Augmented Generation (RAG) workflow using modern LangChain Expression Language- langchain_core LCEL:

- **Semantic Document Indexing (FAISS):** All clinical notes are converted into dense vectors using l3cube-pune/bengali-sentence-bert-nli and stored in an in-memory FAISS vector index. When a user asks a medical question in Bengali then the retriever fetches the top-k most relevant clinical notes in milliseconds.

- **LCEL Execution Graph:** Our chain pipes the retrieved context and user question directly through {context: retriever | format_docs, question: RunnablePassthrough()} | PROMPT | llm | StrOutputParser(), ensuring zero boilerplate overhead and clean string outputs.

4.3 Strict Bengali System Prompting & Anti-Hallucination Guardrails

Large Language Models often try to answer questions using outside pre-trained knowledge or slip into English. In clinical settings, guessing a drug dosage is unacceptable. We engineered our prompt using native Llama-3 instruction control tokens (<|begin_of_text|><|start_header_id|>system<|end_header_id|>) to enforce two ironclad rules:

- **100% Grounded Clinical Answers:** The assistant is instructed to answer strictly from the provided clinical notes (Context). On in-corpus queries such as 'প্যারাসিটামল সাধারণত কী মাত্রায় দেওয়া হয়?', the pipeline accurately retrieves the exact patient notes and outputs the grounded Bengali dosage.
- **Deterministic Refusal on Missing Information:** If the answer is not present in the retrieved notes, the system prompt forbids guessing or lecturing and forces the model to output the exact Bengali phrase: 'তথ্যটি নোটে পাওয়া যায়নি'.

5. Hardware Resource & VRAM Management

To run three heavy deep learning components—a 25-epoch BanglaBERT NER trainer, a 768-dimensional Sentence-BERT embedder and a 8-billion-parameter Llama-3 model inside a single Google Colab notebook is a major system engineering challenge. After 25 epochs of NER training, PyTorch's GPU memory becomes heavily fragmented which leave insufficient contiguous VRAM for Llama-3.

Team Pragyan implemented a four-part memory management strategy to guarantee smooth and crash-free execution from top to bottom:

- **Strategic CPU Offloading for Embeddings & FAISS:** We explicitly pinned the l3cube-pune/bengali-sentence-bert-nli embedding model and the FAISS vector database to the system CPU (model_kwargs={'device': 'cpu'}). Since encoding 762 short clinical notes on CPU takes only a few seconds, this saves nearly 1.5 GB of precious GPU VRAM at zero performance cost.
- **Aggressive Variable Deletion & Garbage Collection:** Immediately before loading Llama-3, our pipeline executes a dedicated cleanup block that checks globals()/locals() and explicitly deletes the BanglaBERT model and trainer objects ‘del trainer, del model’, followed by Python's gc.collect() and PyTorch's torch.cuda.empty_cache(). This flushes leftover optimizer states and fragmented tensors out of the T4 GPU.
- **Persistent Google Drive Checkpoint:** It is because Colab wipes local disk storage when a session disconnects, we exported the trained 25-epoch NER weights and tokenizer permanently to Google Drive. During hidden test set evaluation, the pipeline loads the pre-trained weights in 2 seconds via AutoModelForTokenClassification.from_pretrained() by completely bypassing the 25-epoch training loop.
- **Relative Path Standardization for Submission Packaging:** All dataset reads, checkpoint loads and plot exports are structured so that when packaged into Pragyan_MedBhasha.zip, the evaluation script can run end-to-end without hardcoded path errors.

6. Final Hidden Test Set Evaluation Summary

Pipeline Stage	Validation Benchmark (762 Notes)	Hidden Test Set Result
1. BanglaBERT NER (10 Classes)	Peak F1: 0.8655 | Peak Recall: 0.8977	140742
2. 5D PCA + K-Means (k = 6)	Silhouette: 0.2225 | DBI: 1.3319	Silhouette: 0.2428 | DBI: 1.2412
3. FAISS + 4-Bit Llama-3 RAG	Verified Bengali Grounded Dosage & Guardrail	60 Grounded Bengali Responses Generated (Dual-Corpus JSONL)
4. Peak GPU VRAM Footprint	< 11.5 GB / 16.0 GB T4 GPU (Zero OOM)	Verified Clean Execution